

DNA_Sequence_Classification_G5_New_Data

April 26, 2026

```
[1]: import os
os.environ["CUDA_VISIBLE_DEVICES"] = "-1"

[2]: import os
import random
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.utils.class_weight import compute_class_weight
from sklearn.metrics import (
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    classification_report,
    confusion_matrix,
    roc_curve,
    auc,
    precision_recall_curve,
    average_precision_score
)

from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier

import tensorflow as tf
from tensorflow.keras.models import Model
from tensorflow.keras.layers import (
    Input,
    Conv1D,
    MaxPooling1D,
    BatchNormalization,
    Dropout,
    Bidirectional,
    LSTM,
```

```

        Dense,
        GlobalMaxPooling1D
    )
    from tensorflow.keras.optimizers import Adam
    from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau,
    ↪ModelCheckpoint
    from tensorflow.keras.utils import to_categorical

    SEED = 42

    os.environ["PYTHONHASHSEED"] = str(SEED)
    random.seed(SEED)
    np.random.seed(SEED)
    tf.random.set_seed(SEED)

    print("Libraries imported successfully.")
    print("TensorFlow version:", tf.__version__)
    print("Seed set to:", SEED)

```

Libraries imported successfully.
 TensorFlow version: 2.19.0
 Seed set to: 42

```

[3]: from datasets import load_dataset

dataset = load_dataset("dnagpt/dna_splice_site_prediction")
df = dataset["train"].to_pandas()

print("Dataset loaded successfully.")
print("Shape:", df.shape)
print("Columns:", df.columns.tolist())
display(df.head())

```

```

/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:93:
UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab
(https://huggingface.co/settings/tokens), set it as secret in your Google Colab
and restart your session.
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access
public models or datasets.
    warnings.warn(

README.md:   0%|          | 0.00/314 [00:00<?, ?B/s]

Warning: You are sending unauthenticated requests to the HF Hub. Please set a
HF_TOKEN to enable higher rate limits and faster downloads.
WARNING:huggingface_hub.utils._http:Warning: You are sending unauthenticated

```

requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster downloads.

```
data/train-00000-of-00001.parquet: 0%|          | 0.00/8.80M [00:00<?, ?B/s]
```

```
Generating train split: 0%|          | 0/45619 [00:00<?, ? examples/s]
```

Dataset loaded successfully.

Shape: (45619, 2)

Columns: ['sequence', 'label']

	sequence	label
0	GTCCAGCCAACATGGAGAAGTACATTTCTAGGGACTAAACCTGGCT...	1
1	GAGATTACAGGCATGAGCCACCATACTCAGCCATAGTTTCCTAATC...	0
2	CAGGAGATATCCAGTTAAATTAATTAGCTTTTCTTTCCACTCCAAC...	2
3	GCCCAACTGTAAACATCAGTGCACATGTATTGATGGCGCCGTGGGC...	2
4	GCATACTTGTCTTCAGTATCGTGTGTCTGATGCTGCTGAGCGTGG...	2

```
[4]: print("Columns:", df.columns.tolist())
```

```
SEQ_COL = df.columns[0]
```

```
LABEL_COL = df.columns[1]
```

```
print("Using sequence column:", SEQ_COL)
```

```
print("Using label column:", LABEL_COL)
```

Columns: ['sequence', 'label']

Using sequence column: sequence

Using label column: label

```
[5]: print("Class distribution:")
print(df[LABEL_COL].value_counts().sort_index())
print()
```

```
sequence_lengths = df[SEQ_COL].astype(str).apply(len)
```

```
print("Sequence length statistics:")
```

```
print(sequence_lengths.describe())
```

```
plt.figure(figsize=(8, 5))
```

```
plt.hist(sequence_lengths, bins=30)
```

```
plt.title("Sequence Length Distribution")
```

```
plt.xlabel("Sequence Length")
```

```
plt.ylabel("Frequency")
```

```
plt.grid(True)
```

```
plt.show()
```

Class distribution:

label

0 10000

1 9999


2 25620

Name: count, dtype: int64

Sequence length statistics:

count	45619.0
mean	400.0
std	0.0
min	400.0
25%	400.0
50%	400.0
75%	400.0
max	400.0

Name: sequence, dtype: float64



```
[6]: X_seq = df[SEQ_COL].astype(str).str.upper().values
      y = df[LABEL_COL].values

      print("Number of sequences:", len(X_seq))
      print("Number of labels:", len(y))
      print("Original classes:", sorted(pd.Series(y).unique()))
```

Number of sequences: 45619

Number of labels: 45619

Original classes: [np.int64(0), np.int64(1), np.int64(2)]

```
[7]: label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)

print("Encoded classes:", sorted(np.unique(y_encoded)))

y_categorical = to_categorical(y_encoded)

print("Categorical label shape:", y_categorical.shape)
```

Encoded classes: [np.int64(0), np.int64(1), np.int64(2)]
Categorical label shape: (45619, 3)

```
[8]: MAX_LEN = max(len(seq) for seq in X_seq)

print("Using MAX_LEN =", MAX_LEN)

mapping = {
    "A": [1, 0, 0, 0],
    "C": [0, 1, 0, 0],
    "G": [0, 0, 1, 0],
    "T": [0, 0, 0, 1],
    "N": [0, 0, 0, 0]
}

def encode_sequence(seq, max_len):
    seq = seq.upper()

    if len(seq) < max_len:
        seq = seq + "N" * (max_len - len(seq))
    else:
        seq = seq[:max_len]

    return np.array([mapping.get(base, [0, 0, 0, 0]) for base in seq], dtype=np.
↪float32)

X = np.array([encode_sequence(seq, MAX_LEN) for seq in X_seq])

print("X shape:", X.shape)
print("y shape:", y_categorical.shape)
```

Using MAX_LEN = 400
X shape: (45619, 400, 4)
y shape: (45619, 3)

```
[9]: X_train, X_temp, y_train, y_temp, y_train_raw, y_temp_raw = train_test_split(
    X,
    y_categorical,
    y_encoded,
```

```

        test_size=0.20,
        random_state=SEED,
        stratify=y_encoded
    )

X_val, X_test, y_val, y_test, y_val_raw, y_test_raw = train_test_split(
    X_temp,
    y_temp,
    y_temp_raw,
    test_size=0.50,
    random_state=SEED,
    stratify=y_temp_raw
)

print("Training set:", X_train.shape, y_train.shape)
print("Validation set:", X_val.shape, y_val.shape)
print("Test set:", X_test.shape, y_test.shape)

```

Training set: (36495, 400, 4) (36495, 3)
 Validation set: (4562, 400, 4) (4562, 3)
 Test set: (4562, 400, 4) (4562, 3)

```

[10]: class_weights_values = compute_class_weight(
        class_weight="balanced",
        classes=np.unique(y_train_raw),
        y=y_train_raw
    )

class_weights = dict(enumerate(class_weights_values))

print("Class weights:")
print(class_weights)

```

Class weights:
 {0: np.float64(1.520625), 1: np.float64(1.520815101887736), 2:
 np.float64(0.5935304449648712)}

```

[11]: input_layer = Input(shape=(MAX_LEN, 4))

x = Conv1D(filters=64, kernel_size=5, activation="relu",
padding="same")(input_layer)
x = BatchNormalization()(x)
x = MaxPooling1D(pool_size=2)(x)
x = Dropout(0.3)(x)

x = Conv1D(filters=128, kernel_size=3, activation="relu", padding="same")(x)
x = BatchNormalization()(x)
x = MaxPooling1D(pool_size=2)(x)

```

```

x = Dropout(0.3)(x)

x = Bidirectional(LSTM(64, return_sequences=True))(x)
x = Dropout(0.3)(x)

x = GlobalMaxPooling1D()(x)
x = Dense(64, activation="relu")(x)
x = Dropout(0.3)(x)

output_layer = Dense(3, activation="softmax")(x)

model = Model(inputs=input_layer, outputs=output_layer)

model.summary()

```

Model: "functional"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 400, 4)	0
conv1d (Conv1D)	(None, 400, 64)	1,344
batch_normalization (BatchNormalization)	(None, 400, 64)	256
max_pooling1d (MaxPooling1D)	(None, 200, 64)	0
dropout (Dropout)	(None, 200, 64)	0
conv1d_1 (Conv1D)	(None, 200, 128)	24,704
batch_normalization_1 (BatchNormalization)	(None, 200, 128)	512
max_pooling1d_1 (MaxPooling1D)	(None, 100, 128)	0
dropout_1 (Dropout)	(None, 100, 128)	0
bidirectional (Bidirectional)	(None, 100, 128)	98,816
dropout_2 (Dropout)	(None, 100, 128)	0
global_max_pooling1d (GlobalMaxPooling1D)	(None, 128)	0

dense (Dense)	(None, 64)	8,256
dropout_3 (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 3)	195

Total params: 134,083 (523.76 KB)

Trainable params: 133,699 (522.26 KB)

Non-trainable params: 384 (1.50 KB)

```
[12]: model.compile(
    optimizer=Adam(learning_rate=0.001),
    loss="categorical_crossentropy",
    metrics=["accuracy"]
)

print("Model compiled successfully.")
```

Model compiled successfully.

```
[13]: early_stop = EarlyStopping(
    monitor="val_loss",
    patience=8,
    restore_best_weights=True
)

reduce_lr = ReduceLROnPlateau(
    monitor="val_loss",
    factor=0.5,
    patience=4,
    min_lr=1e-6,
    verbose=1
)

save_best = ModelCheckpoint(
    "best_splice_model.keras",
    monitor="val_loss",
    save_best_only=True,
    verbose=1
)

history = model.fit(
```



```

X_train,
y_train,
validation_data=(X_val, y_val),
epochs=50,
batch_size=32,
class_weight=class_weights,
callbacks=[early_stop, reduce_lr, save_best],
verbose=1
)

```

Epoch 1/50

1141/1141 0s 63ms/step -

accuracy: 0.3170 - loss: 1.1108

Epoch 1: val_loss improved from None to 1.06832, saving model to best_splice_model.keras

Epoch 1: finished saving model to best_splice_model.keras

1141/1141 80s 66ms/step -

accuracy: 0.3293 - loss: 1.1000 - val_accuracy: 0.4459 - val_loss: 1.0683 - learning_rate: 0.0010

Epoch 2/50

1141/1141 0s 62ms/step -

accuracy: 0.3693 - loss: 1.0853

Epoch 2: val_loss improved from 1.06832 to 0.99961, saving model to best_splice_model.keras

Epoch 2: finished saving model to best_splice_model.keras

1141/1141 74s 65ms/step -

accuracy: 0.3872 - loss: 1.0741 - val_accuracy: 0.5449 - val_loss: 0.9996 - learning_rate: 0.0010

Epoch 3/50

1140/1141 0s 62ms/step -

accuracy: 0.5127 - loss: 0.9528

Epoch 3: val_loss improved from 0.99961 to 0.94630, saving model to best_splice_model.keras

Epoch 3: finished saving model to best_splice_model.keras

1141/1141 73s 64ms/step -

accuracy: 0.5253 - loss: 0.9216 - val_accuracy: 0.4811 - val_loss: 0.9463 - learning_rate: 0.0010

Epoch 4/50

1140/1141 0s 62ms/step -

accuracy: 0.5471 - loss: 0.8422

Epoch 4: val_loss improved from 0.94630 to 0.88474, saving model to best_splice_model.keras

Epoch 4: finished saving model to best_splice_model.keras

1141/1141 74s 65ms/step -

```

accuracy: 0.5608 - loss: 0.8129 - val_accuracy: 0.4836 - val_loss: 0.8847 -
learning_rate: 0.0010
Epoch 5/50
1140/1141          0s 63ms/step -
accuracy: 0.6038 - loss: 0.7457
Epoch 5: val_loss improved from 0.88474 to 0.69367, saving model to
best_splice_model.keras

Epoch 5: finished saving model to best_splice_model.keras
1141/1141          75s 66ms/step -
accuracy: 0.6349 - loss: 0.7063 - val_accuracy: 0.6721 - val_loss: 0.6937 -
learning_rate: 0.0010
Epoch 6/50
1140/1141          0s 63ms/step -
accuracy: 0.7238 - loss: 0.5677
Epoch 6: val_loss improved from 0.69367 to 0.56855, saving model to
best_splice_model.keras

Epoch 6: finished saving model to best_splice_model.keras
1141/1141          75s 65ms/step -
accuracy: 0.7429 - loss: 0.5413 - val_accuracy: 0.7409 - val_loss: 0.5686 -
learning_rate: 0.0010
Epoch 7/50
1141/1141          0s 62ms/step -
accuracy: 0.7857 - loss: 0.4749
Epoch 7: val_loss improved from 0.56855 to 0.53570, saving model to
best_splice_model.keras

Epoch 7: finished saving model to best_splice_model.keras
1141/1141          73s 64ms/step -
accuracy: 0.7870 - loss: 0.4668 - val_accuracy: 0.7611 - val_loss: 0.5357 -
learning_rate: 0.0010
Epoch 8/50
1140/1141          0s 61ms/step -
accuracy: 0.8109 - loss: 0.4148
Epoch 8: val_loss did not improve from 0.53570
1141/1141          73s 64ms/step -
accuracy: 0.8082 - loss: 0.4149 - val_accuracy: 0.7337 - val_loss: 0.5794 -
learning_rate: 0.0010
Epoch 9/50
1140/1141          0s 61ms/step -
accuracy: 0.8088 - loss: 0.4112
Epoch 9: val_loss improved from 0.53570 to 0.49941, saving model to
best_splice_model.keras

Epoch 9: finished saving model to best_splice_model.keras
1141/1141          73s 64ms/step -
accuracy: 0.8200 - loss: 0.3920 - val_accuracy: 0.7720 - val_loss: 0.4994 -

```

```

learning_rate: 0.0010
Epoch 10/50
1140/1141          0s 61ms/step -
accuracy: 0.8333 - loss: 0.3676
Epoch 10: val_loss improved from 0.49941 to 0.42970, saving model to
best_splice_model.keras

Epoch 10: finished saving model to best_splice_model.keras
1141/1141          73s 64ms/step -
accuracy: 0.8368 - loss: 0.3633 - val_accuracy: 0.8121 - val_loss: 0.4297 -
learning_rate: 0.0010
Epoch 11/50
1140/1141          0s 62ms/step -
accuracy: 0.8375 - loss: 0.3553
Epoch 11: val_loss improved from 0.42970 to 0.41092, saving model to
best_splice_model.keras

Epoch 11: finished saving model to best_splice_model.keras
1141/1141          74s 64ms/step -
accuracy: 0.8405 - loss: 0.3493 - val_accuracy: 0.8286 - val_loss: 0.4109 -
learning_rate: 0.0010
Epoch 12/50
1140/1141          0s 62ms/step -
accuracy: 0.8469 - loss: 0.3379
Epoch 12: val_loss did not improve from 0.41092
1141/1141          73s 64ms/step -
accuracy: 0.8512 - loss: 0.3313 - val_accuracy: 0.8165 - val_loss: 0.4239 -
learning_rate: 0.0010
Epoch 13/50
1140/1141          0s 62ms/step -
accuracy: 0.8468 - loss: 0.3294
Epoch 13: val_loss improved from 0.41092 to 0.40746, saving model to
best_splice_model.keras

Epoch 13: finished saving model to best_splice_model.keras
1141/1141          73s 64ms/step -
accuracy: 0.8518 - loss: 0.3231 - val_accuracy: 0.8303 - val_loss: 0.4075 -
learning_rate: 0.0010
Epoch 14/50
1141/1141          0s 63ms/step -
accuracy: 0.8531 - loss: 0.3305
Epoch 14: val_loss improved from 0.40746 to 0.34571, saving model to
best_splice_model.keras

Epoch 14: finished saving model to best_splice_model.keras
1141/1141          74s 65ms/step -
accuracy: 0.8585 - loss: 0.3184 - val_accuracy: 0.8683 - val_loss: 0.3457 -
learning_rate: 0.0010

```

Epoch 15/50
1140/1141 0s 62ms/step -
accuracy: 0.8640 - loss: 0.3084
Epoch 15: val_loss did not improve from 0.34571
1141/1141 74s 65ms/step -
accuracy: 0.8626 - loss: 0.3091 - val_accuracy: 0.8665 - val_loss: 0.3490 -
learning_rate: 0.0010
Epoch 16/50
1141/1141 0s 61ms/step -
accuracy: 0.8702 - loss: 0.2995
Epoch 16: val_loss did not improve from 0.34571
1141/1141 73s 64ms/step -
accuracy: 0.8699 - loss: 0.2969 - val_accuracy: 0.8542 - val_loss: 0.3637 -
learning_rate: 0.0010
Epoch 17/50
1140/1141 0s 61ms/step -
accuracy: 0.8669 - loss: 0.3014
Epoch 17: val_loss improved from 0.34571 to 0.32925, saving model to
best_splice_model.keras

Epoch 17: finished saving model to best_splice_model.keras
1141/1141 72s 63ms/step -
accuracy: 0.8702 - loss: 0.2907 - val_accuracy: 0.8726 - val_loss: 0.3292 -
learning_rate: 0.0010
Epoch 18/50
1140/1141 0s 61ms/step -
accuracy: 0.8766 - loss: 0.2824
Epoch 18: val_loss did not improve from 0.32925
1141/1141 72s 63ms/step -
accuracy: 0.8776 - loss: 0.2794 - val_accuracy: 0.8150 - val_loss: 0.4143 -
learning_rate: 0.0010
Epoch 19/50
1141/1141 0s 61ms/step -
accuracy: 0.8776 - loss: 0.2775
Epoch 19: val_loss improved from 0.32925 to 0.32124, saving model to
best_splice_model.keras

Epoch 19: finished saving model to best_splice_model.keras
1141/1141 72s 63ms/step -
accuracy: 0.8770 - loss: 0.2811 - val_accuracy: 0.8742 - val_loss: 0.3212 -
learning_rate: 0.0010
Epoch 20/50
1140/1141 0s 61ms/step -
accuracy: 0.8793 - loss: 0.2703
Epoch 20: val_loss improved from 0.32124 to 0.31234, saving model to
best_splice_model.keras

Epoch 20: finished saving model to best_splice_model.keras

```

1141/1141          72s 63ms/step -
accuracy: 0.8790 - loss: 0.2691 - val_accuracy: 0.8746 - val_loss: 0.3123 -
learning_rate: 0.0010
Epoch 21/50
1140/1141          0s 61ms/step -
accuracy: 0.8745 - loss: 0.2770
Epoch 21: val_loss did not improve from 0.31234
1141/1141          72s 63ms/step -
accuracy: 0.8803 - loss: 0.2680 - val_accuracy: 0.8621 - val_loss: 0.3410 -
learning_rate: 0.0010
Epoch 22/50
1140/1141          0s 61ms/step -
accuracy: 0.8824 - loss: 0.2634
Epoch 22: val_loss improved from 0.31234 to 0.28362, saving model to
best_splice_model.keras

Epoch 22: finished saving model to best_splice_model.keras
1141/1141          72s 63ms/step -
accuracy: 0.8852 - loss: 0.2613 - val_accuracy: 0.8919 - val_loss: 0.2836 -
learning_rate: 0.0010
Epoch 23/50
1140/1141          0s 61ms/step -
accuracy: 0.8857 - loss: 0.2589
Epoch 23: val_loss did not improve from 0.28362
1141/1141          73s 64ms/step -
accuracy: 0.8889 - loss: 0.2529 - val_accuracy: 0.8781 - val_loss: 0.3233 -
learning_rate: 0.0010
Epoch 24/50
1140/1141          0s 62ms/step -
accuracy: 0.8912 - loss: 0.2534
Epoch 24: val_loss did not improve from 0.28362
1141/1141          74s 64ms/step -
accuracy: 0.8933 - loss: 0.2512 - val_accuracy: 0.8889 - val_loss: 0.2968 -
learning_rate: 0.0010
Epoch 25/50
1140/1141          0s 62ms/step -
accuracy: 0.8893 - loss: 0.2521
Epoch 25: val_loss did not improve from 0.28362
1141/1141          73s 64ms/step -
accuracy: 0.8927 - loss: 0.2484 - val_accuracy: 0.8873 - val_loss: 0.3004 -
learning_rate: 0.0010
Epoch 26/50
1141/1141          0s 62ms/step -
accuracy: 0.8915 - loss: 0.2483
Epoch 26: ReduceLROnPlateau reducing learning rate to 0.0005000000237487257.

Epoch 26: val_loss did not improve from 0.28362
1141/1141          73s 64ms/step -

```

```

accuracy: 0.8917 - loss: 0.2467 - val_accuracy: 0.8950 - val_loss: 0.2905 -
learning_rate: 0.0010
Epoch 27/50
1141/1141          0s 62ms/step -
accuracy: 0.9045 - loss: 0.2302
Epoch 27: val_loss did not improve from 0.28362
1141/1141          74s 65ms/step -
accuracy: 0.9069 - loss: 0.2206 - val_accuracy: 0.8801 - val_loss: 0.3270 -
learning_rate: 5.0000e-04
Epoch 28/50
1140/1141          0s 62ms/step -
accuracy: 0.9104 - loss: 0.2128
Epoch 28: val_loss did not improve from 0.28362
1141/1141          73s 64ms/step -
accuracy: 0.9095 - loss: 0.2145 - val_accuracy: 0.8926 - val_loss: 0.2886 -
learning_rate: 5.0000e-04
Epoch 29/50
1140/1141          0s 62ms/step -
accuracy: 0.9073 - loss: 0.2150
Epoch 29: val_loss did not improve from 0.28362
1141/1141          73s 64ms/step -
accuracy: 0.9102 - loss: 0.2088 - val_accuracy: 0.8935 - val_loss: 0.2910 -
learning_rate: 5.0000e-04
Epoch 30/50
1140/1141          0s 62ms/step -
accuracy: 0.9075 - loss: 0.2132
Epoch 30: val_loss improved from 0.28362 to 0.27044, saving model to
best_splice_model.keras

Epoch 30: finished saving model to best_splice_model.keras
1141/1141          74s 64ms/step -
accuracy: 0.9088 - loss: 0.2103 - val_accuracy: 0.9038 - val_loss: 0.2704 -
learning_rate: 5.0000e-04
Epoch 31/50
1141/1141          0s 63ms/step -
accuracy: 0.9125 - loss: 0.2062
Epoch 31: val_loss did not improve from 0.27044
1141/1141          74s 65ms/step -
accuracy: 0.9140 - loss: 0.2016 - val_accuracy: 0.8943 - val_loss: 0.2850 -
learning_rate: 5.0000e-04
Epoch 32/50
1141/1141          0s 62ms/step -
accuracy: 0.9141 - loss: 0.2023
Epoch 32: val_loss did not improve from 0.27044
1141/1141          74s 65ms/step -
accuracy: 0.9156 - loss: 0.1986 - val_accuracy: 0.8952 - val_loss: 0.2887 -
learning_rate: 5.0000e-04
Epoch 33/50

```

```

1140/1141          0s 62ms/step -
accuracy: 0.9145 - loss: 0.2034
Epoch 33: val_loss did not improve from 0.27044
1141/1141          74s 65ms/step -
accuracy: 0.9159 - loss: 0.1984 - val_accuracy: 0.9038 - val_loss: 0.2715 -
learning_rate: 5.0000e-04
Epoch 34/50
1140/1141          0s 62ms/step -
accuracy: 0.9125 - loss: 0.2063
Epoch 34: ReduceLROnPlateau reducing learning rate to 0.0002500000118743628.


Epoch 34: val_loss did not improve from 0.27044
1141/1141          74s 64ms/step -
accuracy: 0.9143 - loss: 0.1993 - val_accuracy: 0.8913 - val_loss: 0.2909 -
learning_rate: 5.0000e-04
Epoch 35/50
1140/1141          0s 62ms/step -
accuracy: 0.9201 - loss: 0.1870
Epoch 35: val_loss did not improve from 0.27044
1141/1141          74s 65ms/step -
accuracy: 0.9220 - loss: 0.1814 - val_accuracy: 0.8897 - val_loss: 0.3009 -
learning_rate: 2.5000e-04
Epoch 36/50
1140/1141          0s 62ms/step -
accuracy: 0.9240 - loss: 0.1841
Epoch 36: val_loss did not improve from 0.27044
1141/1141          73s 64ms/step -
accuracy: 0.9236 - loss: 0.1844 - val_accuracy: 0.8829 - val_loss: 0.3119 -
learning_rate: 2.5000e-04
Epoch 37/50
1141/1141          0s 62ms/step -
accuracy: 0.9237 - loss: 0.1757
Epoch 37: val_loss did not improve from 0.27044
1141/1141          73s 64ms/step -
accuracy: 0.9264 - loss: 0.1723 - val_accuracy: 0.8891 - val_loss: 0.3006 -
learning_rate: 2.5000e-04
Epoch 38/50
1141/1141          0s 62ms/step -
accuracy: 0.9232 - loss: 0.1822
Epoch 38: ReduceLROnPlateau reducing learning rate to 0.0001250000059371814.


Epoch 38: val_loss did not improve from 0.27044
1141/1141          73s 64ms/step -
accuracy: 0.9251 - loss: 0.1755 - val_accuracy: 0.8979 - val_loss: 0.2791 -
learning_rate: 2.5000e-04

```

```
[14]: plt.figure(figsize=(10, 5))
plt.plot(history.history["loss"], label="Training Loss")
plt.plot(history.history["val_loss"], label="Validation Loss")
plt.title("Training vs Validation Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.grid(True)
plt.show()

plt.figure(figsize=(10, 5))
plt.plot(history.history["accuracy"], label="Training Accuracy")
plt.plot(history.history["val_accuracy"], label="Validation Accuracy")
plt.title("Training vs Validation Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.legend()
plt.grid(True)
plt.show()
```





```
[15]: y_pred_prob = model.predict(X_test)
y_pred = np.argmax(y_pred_prob, axis=1)
y_true = np.argmax(y_test, axis=1)

accuracy = accuracy_score(y_true, y_pred)
precision = precision_score(y_true, y_pred, average="macro")
recall = recall_score(y_true, y_pred, average="macro")
f1 = f1_score(y_true, y_pred, average="macro")


print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)

print("\nClassification Report:\n")
print(classification_report(y_true, y_pred))
```

```
143/143          3s 21ms/step
Accuracy: 0.9050854888206927
Precision: 0.8842133991206605
Recall: 0.9276088992974239
F1 Score: 0.9021093718610813
```

Classification Report:

precision	recall	f1-score	support
-----------	--------	----------	---------



```
[25]: model.save("final_cnn_bilstm_splice_model.keras")
      results_df.to_csv("model_comparison_results.csv", index=False)
      report_df.to_csv("classification_report_table.csv")

      with open("final_model_metrics.txt", "w") as f:
          f.write(f"Accuracy: {accuracy}\n")
          f.write(f"Precision: {precision}\n")
          f.write(f"Recall: {recall}\n")
          f.write(f"F1 Score: {f1}\n")
          f.write(f"AUC Scores: {roc_auc}\n")
          f.write(f"Average Precision Scores: {avg_precision}\n")

      print("Model and results saved successfully.")
```

Model and results saved successfully.